

IMPERIAL COLLEGE LONDON

DEPARTMENT OF COMPUTING

Symbolic Reinforcement Learning using Inductive Logic Programming

Author:
Kiyohito Kunii

Supervisor:
Professor Alessandra Russo

Submitted in partial fulfillment of the requirements for the MSc degree in MSc in
Computing Science of Imperial College London

June 2018

Contents

1	Introduction	2
2	Background	3
2.1	Inductive Logic Programming	3
2.1.1	Stable Model Semantics	3
2.1.2	Answer Set Programming (ASP)	3
2.1.3	Learning from Answer Sets (LAS)	5
2.1.4	Inductive Learning of Answer Set Programs (ILASP)	5
2.2	Reinforcement Learning	5
2.2.1	Markov Decision Process(MDP)	5
2.2.2	Temporal-Difference (TD) Learning	6
2.2.3	Q-Learning	7
2.3	Deep Reinforcement Learning	8
2.3.1	Convolutional Neural Networks (CNN)	8
2.3.2	Deep Q-Networks (DQNs)	8
2.4	GVGAI Framework	8
3	Related Work	9
3.1	Deep Reinforcement Learning and its issues	9
3.1.1	Data efficient learning	9
3.1.2	Transfer learning	9
3.1.3	Planning	9
3.1.4	Transparency	9
3.2	Symbolic Reinforcement Learning	10
4	Project Overview	11
4.1	Motivation	11
4.2	Objectives	11
4.3	Project outline	11
4.4	Contribution	12
4.5	Methods	12
5	Ethics Checklist	13
6	References	16

Chapter 1

Introduction

There has been successful applications of deep reinforcement learning (DRL) a number of domains, such as games and robotics (TODO INSERT REFERENCES). DRL is considered to be a step towards artificial general intelligence. As pointed by XXX, however, there are still a number of issues to overcome in this method. First, it requires a large amount of data for training the model, which requires a long time of learning process. Second, it is considered to be a black-box, meaning the decision making process is unknown to human user and therefore lacks explanation ability. Third there is no thoughts process to the decision making, which, as XX points out, is a fundamental to the artificial general intelligence. To tackle these problems, there are 3 main streams of research on this field. First main research is focused on applying Bayesian statistics XXX, the second main research is XXX. Recently, the XXX attempted to incorporate symbolic representations into the system to achieve more data-efficient learning, which shows a promising results of this approach.

In this paper, I further explore incorporation of symbolic machine learning into reinforcement learning to achieve data-efficient learning using Inductive Learning of Answer Set Programs (ILASP), which is the state-of-art symbolic learning method that can be applied to incomplete and more complex environment. This research is inspired by [1], but in this paper I explore symbolic representations process and include more learning aspect.

This background report is organised as follows: Chapter 2 introduces

Chapter 2

Background

This section introduces introduces Inductive Logic Programming (Section 2.1), Reinforcement Learning (Section 2.2), Deep Reinforcement Learning (Section 2.3) and GVGA Framework, an open-source game platform for AI methods (Section 2.4), which provide the foundations of this work.

2.1 Inductive Logic Programming

Inductive Logic Programming (ILP) is a subfield of research area aimed at the intersection of machine learning and logic programming (MUGGLETON 1991). The purpose of ILP is to derive a hypothesis H that is a solution of a learning task, which covers all positive examples and any of negative examples.

Limitations of Inductive Logic Programming:

$$B \cap H \models E \quad (2.1)$$

TODO Herbrand Model

TODO Least Herbrand Model

Definite Logic Program is a set of definite rule where:

a definite rule is of the form $h \leftarrow a_1, \dots, a_n$, where h, a_1, \dots, a_n are all atoms.
whereas

Normal Logic Program is a set of normal rule where

a normal rule is of the form $h \leftarrow a_1, \dots, a_n, \text{not } b_1, \dots, \text{not } b_n$ where h is the head of the rule, and $a_1, \dots, a_n, b_1, \dots, b_n$ are the body of the rule (both the head and body are all atoms).

2.1.1 Stable Model Semantics

Stable Model of a normal logic program

2.1.2 Answer Set Programming (ASP)

Answer Set Programming (ASP) is a declarative programming

Literals

Negation as a failure

constrains is of the form

Clingo

$\leftarrow a_1, \dots, a_n, notb_1, \dots, notb_n$

Constraints are to filtering any irrelevant answer sets.

Syntax Examples

There are two types of constraints: soft and hard constraints. Soft constraint is XXX

Hard constraint is XXX

optimisation statement is of the form.

Which is useful to order the answer sets in terms of preference.

Syntax examples.

Aggregate

choice rules

An ASP program P is normal logic program with addition of choice rules, constraints and optimisation statement.

Answer set of P is

Cautious Induction

Sakama 2008

no concept of negative examples in this paper.

Cautious Induction task is of the form $\langle B, E^+, E^- \rangle$ where:

B is the background knowledge

E^+ is a set of positive examples

E^- is a set of negative examples

$H \in ILP_{\text{cautious}} \langle B, E^+, E^- \rangle$ if and only if

there is at least one answer set A of $B \cup H$ such that:

for every anser set A of $B \cup H$: $\forall e \in E^+ : e \in A$

$\forall e \in E^- : e \notin A$

One of the limitation of cautious indction is that XX

Sakama 2009

Brave Induction

Similarly, Brave Induction task is of the form $\langle B, E^+, E^- \rangle$ where:

B is the background knowledge

E^+ is a set of positive examples

E^- is a set of negative examples

$H \in ILP_{\text{brave}} \langle B, E^+, E^- \rangle$ if and only if there is at least one answer set A of $B \cup H$ such that:

$\forall e \in E^+ : e \in A$

$$\forall e \in E^- : e \notin A$$

One of the limitation of brave induction is that it cannot learn constraints as shown in the example

Sakama 2009

2.1.3 Learning from Answer Sets (LAS)

Learning from Answer Sets was developed in [Law et al 2014] to overcome limitations of cautious induction and brave induction.

Examples used in LAS is converted from $\langle E^+, E^- \rangle$ into

Partial Interpretations are of the form $\langle E^{\text{inc}}, E^{\text{exc}} \rangle$

A Herbrand Interpretations extends a partial interpretation if it include all of the inclusions and none of the exclusions.

2.1.4 Inductive Learning of Answer Set Programs (ILASP)

ILASP is an algorithm that is capable of solving LAS tasks

It is based on two fundamental concepts: positive solutions and violating solutions.

A hypothesis H is a positive solution if and only if 1. $H \subseteq S_M$

$$2. \forall e^+ \in E^+ \exists A \in AS(B \cup H) \text{subject to } A \text{ extend } e^+$$

A hypothesis H is a violating solution if and only if 1. $H \subseteq S_M$

$$2. \forall e^+ \in E^+ \exists A \in AS(B \cup H) \text{subject to } A \text{ extend } e^+$$

$$3. \exists e^- \in E^- \exists A \in AS(B \cup H) \text{subject to } A \text{ extend } e^-$$

ILP_{LAS} is positive solutions that are not violating solutions.

ILASP task containing a contex-dependent example

TODO Explain how symbolic learning works

TODO What would you learn in my context? Relationship of the objects? Objects, types, locations and interactions.

2.2 Reinforcement Learning

On-Policy learn policy π from its past experience Off-policy

2.2.1 Markov Decision Process(MDP)

There is an agent who interacts with an environment (Figure 2.1). At each time step, an action taken by the agent affect the environment state and the reward (or penalty) it receives from the action outcome. (TODO Observations) When an agent must make a sequence of decision, the sequential decision problem can be formalised using Markov decision process (MDP). MDPs formally represent a fully observable environment of an agent for reinforcement learning.




Figure 2.1: Agent and Environment

A MDP is of the form $\langle S, A, T_a, R_a, \gamma \rangle$ where:

- S is the set of finite states that is observable in the environment
- A is the set of finite actions taken by the agent
- $T_a(s, s')$ is a state transition in the form of probability matrix $\Pr(S_{t+1} = s' \mid s_t = s, a_t = a)$, which is the probability that action a in state s at time t will result in state s' at time $t+1$.
- R is a reward function $R_a(s, s') = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$, the expected immediate reward that action a in state s at time t will return
- γ is a discount factor $\gamma \in [0, 1]$, which represents the preference of the agent for present rewards over future rewards

In MDPs, there is an element of delayed reward tradeoff between immediate rewards and its delayed reward

A state S_t is Markov if and only if $P[S_{t+1} \mid S_t] = P[S_{t+1} \mid S_1, \dots, S_t]$ therefore the probability of reaching S_{t+1} depends solely on S_t , which captures all the relevant information from earlier history.

A solution to the sequential decision problem is called a policy π , which is a distribution over actions given states.

An optimal policy π^* is the one that maximise the total rewards in the environment. The total reward with a discount factor is

The existence of the discount factor can be justified as follows:

TODO $\mathbb{E}[R_{t+1} \mid S_t = s]$

Reinforcement learning is a method to get approximated optimal solution.

2.2.2 Temporal-Difference (TD) Learning

TD learns directly from episodes of experiences, which can be incomplete. TD does not require knowledge of MDP transitions and rewards (model-free) Sutton 1988

Update value

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t)) \quad (2.2)$$

where $R_{t+1} + \gamma V(S_{t+1})$ is the estimated return (a.k.a TD target)

$R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is TD error.

TD updates the estimate by using the estimates of XXX (bootstrap).

The advantages of TD methods - does not require any model of an environment. - online learning

2.2.3 Q-Learning

model-free learning, where the agent will not rely on the models of the environment.

Q-learning approximate the $Q(a,s)$ from the samples of

the agent only knows about the possible states and actions but the transition state and reward probability functions are unknown.

Q function is the state-action pair

the optimal Q-function $Q^*(s,a)$ represents XXX for the agent to selection action a given that it is in state s .

Q-learning is off-policy TD learning defined in [Watkin 1989], which is of the form:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_a Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$

(2.3)

where α is the learning rate, γ is a discount rate between 0 and 1.

this equation is used to upate action-value function

Model free learning: derectly derive an optimal policy by interacting with the environment without the model

Model free can be done using Monte Carlo Policy evaluation

One way to solve the Bellman Optimality equation is Q-leraning

$U(s) = \max_a Q(s,a)$

a function $Q(S,A)$ which predicts the best action A in state S to maximise the total cumulative rewards.

The function is estimated by Q-learning, which repeately updates $Q(s,a)$ using the Bellman Equation.

TODO INSERT Q-learning ALGORITHM HERE

Epsilon greedy

2.3 Deep Reinforcement Learning

2.3.1 Convolutional Neural Networks (CNN)

2.3.2 Deep Q-Networks (DQNs)

Explain the paper

2.4 GVGAI Framework

The Video Game Definition Language (VGDL)

SpriteSet defines all the sprites available in the game. LevelMapping defines relationships among characters, and the sprites available. InteractionSet defines what events occur when two sprites collide TerminationSet defines the end conditions of the game, and decides whether the player wins or not.

Chapter 3

Related Work

3.1 Deep Reinforcement Learning and its issues

3.1.1 Data efficient learning

3.1.2 Transfer learning

3.1.3 Planning

3.1.4 Transparency

Transparency and interpretable capability of the model is another important aspect for machine learning applications.

XXX[Programmati...] developed a programmatically interpretable reinforcement learning which finds a policy that can be represented in a human-readable programming language.

Two most studied approach for using previous learning exprience is meta-learning and transfer learning

Artificial General Intelligence

Definision of AGI

The history of data-efficient learning What other people have done in this.

The advance of statistical machine learning methods, especially deeep reinforcement learning

AlphaGo, and AlphaGo Zero

Study of symbolic machine learning roots from

Baysian Optimisation

RNN approach

Symbolic Deep reinforcement learning

Some implementation: German paper

The paper was the application of symbolic representatiojns into a very simple game to demonstrate this proof of concept would actually work. By contract, there is active reserach in symbolic machine learning, which focuses on logic=based learning rather than statistical machine learning. For example, XX shows the agent can learn XXX from a noisy examples, with only a very few training examples.

3.2 Symbolic Reinforcement Learning

Incorporation of logic into reinforcement learning dates back to the study of relational reinforcement learning,

More recently there has been a number of attempts to incorporate ASP into reinforcement learning.

There are a number of reseached conducted in applying DNN to symbolic reasoning.

For example,

[From GamePlay to Symbolic Reasoning]

However, as XX points out, 1 pixel of the game could influence the decision making of the agent.

Chapter 4

Project Overview

4.1 Motivation

Why symbolic reinforcement learning is good attempt

Reason 1: Comprehensive by humans → Explainable rather than black-box

Reason 2: Similar to the human, it uses reasoning

Use of previous experience (background knowledge) As discussed in XX, there are many research that attempted to incorporate symbolic reasoning into DRL, but there is not much research on incorporating symbolic learning with reinforcement learning. However there is a room for exploration on this field.

Reason 3: Recent advance of ILASP is promising

Because of the recent advancement of logic-based learning and deep reinforcement learning, combination of both approach would be a next exploration toward artificial general intelligence.

4.2 Objectives

combining the two novel approaches to overcome the problems of the QND

4.3 Project outline

The project outline

Implement baseline performance (DQN, Q-learning, ASPRL)

Due to the complex environment of the chosen game, I will not implement the original DSRL, since the feature extractions from the pixel would only work in a very simple pixel environment,

Pipeline

The basic architecture follows the similar method in XXX, but I also add symbolic learning to these extracted symbolic features using ILASP.

Apply ILASP to the ASP, which involves development of the pipeline of ILASP in Python

Finally use Q-learning that allow

which measurement would you use? (grid world, something else? GVGAL games)
Summarise different types of knowledge representations (Objects ?? relationship?)
Common sense
Implement based on Towards Deep Symbolic Reinforcement Learning
Lastly, I will also test the capability of transfer learning for this new method.

4.4 Contribution

To my knowledge, this is the first time that both symbolic learning method is incorporated into a reinforcement learning to facilitate learning process

4.5 Methods

Chapter 5

Ethics Checklist

	Yes	No
Section 1: HUMAN EMBRYOS/FOETUSES		
Does your project involve Human Embryonic Stem Cells?		✓
Does your project involve the use of human embryos?		✓
Does your project involve the use of human foetal tissues / cells?		✓
Section 2: HUMANS		
Does your project involve human participants?		✓
Section 3: HUMAN CELLS / TISSUES		
Does your project involve human cells or tissues? (Other than from Human Embryos/Foetuses i.e. Section 1)?		✓
Section 4: PROTECTION OF PERSONAL DATA		
Does your project involve personal data collection and/or processing?		✓
Does it involve the collection and/or processing of sensitive personal data (e.g. health, sexual lifestyle, ethnicity, political opinion, religious or philosophical conviction)?		✓
Does it involve processing of genetic information?		✓
Does it involve tracking or observation of participants? It should be noted that this issue is not limited to surveillance or localization data. It also applies to Wan data such as IP address, MACs, cookies etc.		✓
Does your project involve further processing of previously collected personal data (secondary use)? For example Does your project involve merging existing data sets?		✓
Section 5: ANIMALS		
Does your project involve animals?		✓
Section 6: DEVELOPING COUNTRIES		
Does your project involve developing countries?		✓

If your project involves low and/or lower-middle income countries, are any benefit-sharing actions planned?		✓
Could the situation in the country put the individuals taking part in the project at risk?		✓
Section 7: ENVIRONMENTAL PROTECTION AND SAFETY		
Does your project involve the use of elements that may cause harm to the environment, animals or plants?		✓
Does your project deal with endangered fauna and/or flora /protected areas?		✓
Does your project involve the use of elements that may cause harm to humans, including project staff?		✓
Does your project involve other harmful materials or equipment, e.g. high-powered laser systems?		✓
Section 8: DUAL USE		
Does your project have the potential for military applications?		✓
Does your project have an exclusive civilian application focus?		✓
Will your project use or produce goods or information that will require export licenses in accordance with legislation on dual use items?		✓
Does your project affect current standards in military ethics e.g., global ban on weapons of mass destruction, issues of proportionality, discrimination of combatants and accountability in drone and autonomous robotics developments, incendiary or laser weapons?		✓
Section 9: MISUSE		
Does your project have the potential for malevolent/criminal/terrorist abuse?		✓
Does your project involve information on/or the use of biological-, chemical-, nuclear/radiological-security sensitive materials and explosives, and means of their delivery?		✓
Does your project involve the development of technologies or the creation of information that could have severe negative impacts on human rights standards (e.g. privacy, stigmatization, discrimination), if misapplied?	✓	
Does your project have the potential for terrorist or criminal abuse e.g. infrastructural vulnerability studies, cybersecurity related project?		✓
Section 10: LEGAL ISSUES		
Will your project use or produce software for which there are copyright licensing implications?	✓	

Will your project use or produce goods or information for which there are data protection, or other legal implications?		✓
Section 11: OTHER ETHICS ISSUES		
Are there any other ethics issues that should be taken into consideration?	✓	

Chapter 6

References

Bibliography

- [1] Marta Garnelo, Kai Arulkumaran, and Murray Shanahan. Towards Deep Symbolic Reinforcement Learning. sep 2016. pages 2