

IMPERIAL COLLEGE LONDON

DEPARTMENT OF COMPUTING

Symbolic Reinforcement Learning using Inductive Logic Programming (DRAFT)

Author:
Kiyohito Kunii

Supervisor:
Professor Alessandra Russo

Submitted in partial fulfillment of the requirements for the MSc degree in MSc in
Computing Science of Imperial College London

June 2018

Contents

1	Introduction	2
2	Background	4
2.1	Inductive Logic Programming (ILP)	4
2.1.1	Logic Basics	4
2.1.2	Stable Model Semantics	5
2.1.3	Answer Set Programming (ASP) Syntax	6
2.1.4	ILP Under Answer Set Semantics	7
2.1.5	Inductive Learning of Answer Set Programs (ILASP)	8
2.2	Reinforcement Learning (RL)	11
2.2.1	Markov Decision Process (MDP)	11
2.2.2	Policies and Value Functions	12
2.2.3	Model-based vs Model-free Learning	12
2.2.4	Temporal-Difference (TD) Learning	13
2.2.5	Q-Learning	13
3	Related Work	15
4	Project Overview	17
4.1	Proposed Architecture	17
4.1.1	ASP Representation Generation	18
4.1.2	Model Generation and Update using ILASP	18
4.1.3	Generation of Simulated Experience	19
4.2	Further Research and Plan	19
4.2.1	Further Research	19
4.2.2	Experiment Implementation	20
5	Ethics Checklist	21

Chapter 1

Introduction

There has been successful applications of deep reinforcement learning (DRL) in a number of domains, such as video games ([16]), game of Go ([22]) and robotics ([15]). However, there are still a number of issues to overcome in this method. First, it requires a large amount of data for training the model, which requires a long time of learning process. Second, it is considered to be a black-box, meaning that the decision making process is unknown to human user and therefore lacks explanation ability. Third there is no thoughts process to the decision making, such as understanding relational representations or planning. To tackle these problems, several different approaches the researchers have been explored. One of the methods is to incorporate symbolic representations into the system to achieve more data-efficient learning ([8]), which shows a promising potential of this approach. In this paper, we extend this approach and explore a potential of incorporating inductive logic programming into reinforcement learning to solve the above issues. The advantages of symbolic reinforcement learning are, first of all, the decision making mechanism is understandable by humans rather than being black-box. Second, it resembles how human reasons. Like reinforcement learning, there are some aspects of try-and-error in human learning, but humans exploit reasonings to efficiently learn the surrounding or situations. They also effectively use previous experience (background knowledge) when encountered a similar situation. Finally, the recent advance of ILP research enabled us to apply ILP in more complex situations and there are a number of new algorithms that work well in non-monotonic scenarios based on Answer Set Programmings (ASPs).

After [8], especially, there are several researches that further explored incorporation of symbolic reasoning into DRL, but the combining inductive logic programming and reinforcement learning has not been explored. Because of the recent advancement of symbolic machine learning and deep reinforcement learning, combination of both approach would be a next exploration.

In this paper, our objective is to explore this new research field and see how this combination of the two different learning methods could enhance the learning capability. In this paper, I further explore incorporation of symbolic machine learning into reinforcement learning to achieve data-efficient learning using Inductive Learning of Answer Set Programs (ILASP), which is the state-of-art symbolic learning method that can be applied to incomplete and more complex environment.

This background report will be the part of the final report and is organised as follows: In Chapter 2, necessary background of logic, inductive logic programming and reinforcement learning are described for this paper. Chapter 3 discusses previous research in this particular approach and see the drawbacks of each approach. Chapter 4 shows the tentative architecture of our new approach, using Inductive Learning of Answer Set Programs (ILASP) to generate a model of the environment. In Chapter 4.2, we discuss some of the issues we currently face and plan for the implementation . Finally the ethics checklist is provided in Chapter 5.

Chapter 2

Background

This section introduces necessary background of Inductive Logic Programming (Section 2.1) and Reinforcement Learning (Section 2.2), which provide the foundations of this work.

2.1 Inductive Logic Programming (ILP)

Inductive Logic Programming (ILP) is a subfield of machine learning research area aimed at the intersection between machine learning and logic-based knowledge representation [17]. The purpose of ILP is to inductively derive a hypothesis H that is a solution of a learning task, which covers all positive examples and any of negative examples, given a hypothesis language for search space and cover relation ([5]). The possible explanation of covers relation is described by the entailment, as shown in Equation 2.1.

$$B \cap H \models E \quad (2.1)$$

where E consists of positive examples (E^+) and none of the negative examples (E^-). One of the advantage of ILP over statistical machine learning is that the hypothesis the agent learnt can be easily understood by a human, making the decision more transparent rather than black box.

TODO Limitations of Inductive Logic Programming:

In this section, we briefly introduce basic logic notions, Answer set programming (ASP)

2.1.1 Logic Basics

TODO Satisfiable

To compute an inference task, the syntax of the predicate logic program needs to be converted into Conjunctive Normal Form (CNF), or clausal theory, which consists of a conjunction of clauses, where a clause is disjunction of literals, and a literal can include positive literals and negative literals. A literal is either an atom p or $not\ p$ (negation by failure).

example 2.1.1. Conjunctive Normal Form (Clausal Theory)

$$(rain \vee umbrella \vee rain_shoes) \wedge (sunny \vee roofer)$$

A *Horn clause* is a subset of CNF, which is a clause with at most one positive literal, where a *definite clause* (also called a *rule* or a *fact*) contains exactly one positive literal, and a *denial* (or a *constraint*) is a clause with no positive literals. A *normal clause* extends Horn clause with negation as failure.

2.1.2 Stable Model Semantics

Having defined the syntax of clausal logic, we now introduce its semantics under the context of stable model. The semantics of the logic is based on the notion of *interpretation*, which is defined under a *domain*. A *domain* contains all the objects that exist. For the convenient reasons, we focus on a special interpretations called *Herbrand interpretations* rather than generation interpretations.

The *Herbrand domain* (a.k.a *Herbrand universe*) of clause sets Th is the set of all ground terms that are constants and function symbols appeared in Th .

The *Herbrand base* of Th is the set of all ground predicates that are formed by predicate symbols in Th and terms in the *Herbrand domain*.

The *Herbrand interpretation* of a set of definite clauses Th is a subset of the Herbrand base of Th , which is a set of ground atoms that are true in terms of interpretation.

Interpretation evaluate it to true Interpretation evaluate it to false

A *Herbrand Model* is a Herbrand interpretation if and only if a set Th of clauses is satisfiable. In other words, a set of clauses Th is unsatisfiable if no Herbrand model was found.

The *Herbrand Model* is a minimum Herbrand model if and only if none of its subsets is an Herbrand model. For definite logic programs, there is a unique minimal Herbrand model (the *Least Herbrand Model* $M(P)$). For normal logic programs, there may not be a least Herbrand Model.

Definite Logic Program is a set of definite rule, and a definite rule is of the form $h \leftarrow a_1, \dots, a_n$, where h, a_1, \dots, a_n are all atoms. h is the *head* of the rule and a_1, \dots, a_n are the body of the rule.

Normal Logic Program is a set of normal rule, and a normal rule is of the form $h \leftarrow a_1, \dots, a_n, \text{not } b_1, \dots, \text{not } b_n$ where h is the head of the rule, and $a_1, \dots, a_n, b_1, \dots, b_n$ are the body of the rule (both the head and body are all atoms).

To solve a normal logic program Th , the program needs to be grounded. The *grounding* of Th is the set of all clauses that are $c \in Th$ and variables are replaced by terms in the *Herbrand Domain*.

example 2.1.2. (Grounding)

$$P = \begin{cases} p \leftarrow \text{not } q. \\ q \leftarrow p. \end{cases}$$

ground(P) in Example 2.1.2 is $p:- \text{not } q.$ and $q:- p.$ The algorithm of grounding start with the empty program $Q = \{ \}$ and the relevant grounding is constructed

by adding to each rule R to Q given that R is a ground instance of a rule in P and their positive body literals already occurs in the in the of rules in Q . The algorithm terminates when no more rules can be added to Q . Not only the entire program needs to be grounded in order for ASP solver to work, and, unlike Prolog, but also each rule must be *safe*. A rule R is safe if every variable that occurs in the head of the rule occurs at least once in $\text{body}^+(R)$.

Since there is no unique least Herbrand Model for a normal logic program, [10] defined Stable Model of a normal logic program. In order to obtain the Stable model of P , P needs to be converted using *Reduct* with respect to an interpretation X . First, if the body of any rule in P contains an atom which is not in X , those rules need to be removed. Second, all default negation atoms in the remaining rules in P need to be removed.

example 2.1.3. (Reduct)

$$X = p$$

$$P = \begin{cases} p \leftarrow \text{not } q. \\ q \leftarrow \text{not } p. \\ r \leftarrow p. \end{cases}$$

In the Example 2.1.2, P^X is $p, r \leftarrow p$. A stable model of P is an interpretation X if and only if X is the unique least Herbrand Model of $\text{ground}(P)^X$ in the logic programs.

2.1.3 Answer Set Programming (ASP) Syntax

Answer set of normal logic program P is a stable model, and Answer Set Programming (ASP) is a normal logic program with extensions: constraints, choice rules and optimisation. ASP program consists of a set of rules, where each rule consists of an atom and literals. A literal is either an atom p or its default negation (negation as a failure) $\text{not } p$.

A Constraint of the program P is of the form $\leftarrow a_1, \dots, a_n, \text{not } b_1, \dots, \text{not } b_n$, where the rule has an empty head. The constraint filters any irrelevant answer sets. When computing $\text{ground}(P)_X$, the empty head becomes \perp , which cannot be in the Answer Set. There are two types of constraints: *hard constraints* and *hard constraints*. Hard constraints are strictly satisfied, whereas soft constraints are must not be satisfied but the sum of the violations should be minimised when solving ASP.

Choice rule can express possible outcomes given an action choice, which is of the form $l\{h_1, \dots, h_m\}u \leftarrow a_1, \dots, a_n, \text{not } b_1, \dots, \text{not } b_n$ where l and u are integers and h_i for $1 \leq i \leq m$ are atoms. The head is called *aggregates*.

optimisation statement is useful to order the answer sets in terms of preference, which is of the form $\# \text{minimize}[a_1=w_1, \dots, a_n=w_n]$ or $\# \text{maximize}[a_1=w_1, \dots, a_n=w_n]$ where w_1, \dots, w_n is an integer weight and a_1, \dots, a_n is an ground atom. ASP solvers compute the scores of the weighted sum of the sets of ground atoms based on the true answer sets, and find optimal Answer Sets which either maximise or minimise the score.

Clingo is one of ASP modern solvers that executes the ASP program and returns answer sets of the programs ([9]), and we will use *Clingo* for the implementation of this research.

TODO ASP has true, false and unknown

2.1.4 ILP Under Answer Set Semantics

There are several ILP non-monotonic learning frameworks under the Answer set semantics. We introduce two of them: *Cautious Induction* and *Brave Induction* ([20]), which are foundations of *Learning from Answer Sets* discussed in Section ?? (for other ILP frameworks, see [18], [11], [4] and [5]).

Cautious Induction

Cautious Induction task is of the form $\langle B, E^+, E^- \rangle$, where B is the background knowledge, E^+ is a set of positive examples and E^- is a set of negative examples.

$H \in \text{ILP}_{\text{cautious}} \langle B, E^+, E^- \rangle$ if and only if there is at least one answer set A of $B \cup H$ ($B \cup H$ is satisfiable) such that for every answer set A of $B \cup H$:

$\forall e \in E^+ : e \in A$

$\forall e \in E^- : e \notin A$

example 2.1.4. Cautious Induction (Example)

$$B = \begin{cases} \text{exercises} \leftarrow \text{not eat_out.} \\ \text{eat_out} \leftarrow \text{exercises.} \end{cases}$$

$E_+ = \text{tennis}$

$E_- = \text{eat_out}$

One possible $H \in \text{ILP}_{\text{cautious}}$ is $\{\text{tennis} \leftarrow \text{exercises}, \leftarrow \text{not tennis}\}$.

The limitation of Cautious Induction is that positive examples must be true for all Answer Sets and negative examples must not be included in any of the Answer Sets. These conditions may be too strict in some cases, and Cautious Induction is not able to accept the case where positive examples are true in some of the Answer Sets but not all Answer Sets of the program.

example 2.1.5. Limitation of Cautious Induction (Example)

$1\{\text{situation}(P, \text{awake}), \text{situation}(P, \text{sleep})\}1 :- \text{person}(P).$
 $\text{person}(\text{john}).$

In the example 2.1.4, neither of $\text{situation}(\text{john}, \text{awake})$ or $\text{situation}(\text{john}, \text{sleep})$ is false in all answer sets. In this example, it only return $\text{person}(\text{john})$. Thus no examples could be given to learn the choice rule.

Brave Induction

Brave Induction task is of the form $\langle B, E^+, E^- \rangle$ where, B is the background knowledge, E^+ is a set of positive examples and E^- is a set of negative examples. $H \in$

$ILP_{\text{brave}} \langle B, E^+, E^- \rangle$ if and only if there is at least one answer set A of $B \cup H$ such that:

$$\forall e \in E^+ : e \in A$$

$$\forall e \in E^- : e \notin A$$

example 2.1.6. Brave Induction (Example)

$$B = \begin{cases} \text{exercises} \leftarrow \text{not eat_out.} \\ \text{tennis} \leftarrow \text{holiday} \end{cases}$$

$$E_+ = \text{tennis}$$

$$E_- = \text{eat_out}$$

One possible $H \in ILP_{\text{brave}}$ is $\{\text{tennis}\}$, which returns $\{\text{tennis, holiday, exercises}\}$ as Answer Sets.

The limitation of Brave induction that it cannot learn constraints as shown in the example 2.1.4.

TODO MORE EXPLANATION

example 2.1.7. Limitation of Brave Induction (Example)

$$B = \begin{cases} 1\{\text{situation}(P, \text{awake}), \text{situation}(P, \text{sleep})\} 1 \leftarrow \text{person}(P). \\ \text{person}(C) \leftarrow \text{super_person}(C). \\ \text{super_person}(\text{john}). \end{cases}$$

In order to learn the constraint hypothesis $H = \{ \leftarrow \text{not situation}(P, \text{awake}), \text{super_person}(P) \}$, it is not possible to find an optimal solution.

2.1.5 Inductive Learning of Answer Set Programs (ILASP)

Learning from Answer Sets (LAS)

Learning from Answer Sets (LAS) was developed in [12] to facilitate more complex learning task that neither Cautious Induction nor Brave Induction could learn. Examples used in LAS are *Partial Interpretations*, which are of the form $\langle e^{\text{inc}}, e^{\text{exc}} \rangle$. (called *inclusions* and *exclusions* of e respectively). A Herbrand Interpretations extends a partial interpretation if it include all of e^{inc} and none of e^{exc} .

LAS is of the form $\langle B, S_M, E^+, E^- \rangle$, where B is background knowledge, S_M is hypothesis space, and E^+ and E^- are examples of positive and negative partial interpretations. S_M consists of a set of normal rules, choice rules and constraints. S_M is specified by *language bias* of the learning task using *mode declaration*. Mode declaration, and specifies what can occur in an hypothesis by specifying the predicate, and has two parts: *modeh* and *modeb*. *modeh* and *modeb* are the predicates that can occur in the head of the rule and body of the rule respectively. Language bias is the specification of the language in the hypothesis in order to reduce search space for the hypothesis. Given a learning task T , the set of all possible inductive solutions of T is denoted as $ILP_{\text{LAS}}(T)$, and a hypothesis H is an inductive solution of $ILP_{\text{LAS}}(T) \langle B, S_M, E^+, E^- \rangle$ such that

1. $H \subseteq S_M$
2. $\forall e \in E^+ : \exists A \in \text{Answer Sets}(B \cup H)$ such that A extends e
3. $\forall e \in E^- : \nexists A \in \text{Answer Sets}(B \cup H)$ such that A extends e

example 2.1.8. LAS (Example)

TODO

Inductive Learning of Answer Set Programs (ILASP)

ILASP is an algorithm that is capable of solving LAS tasks, and is based on two fundamental concepts: positive solutions and violating solutions.

A hypothesis H is a positive solution if and only if

1. $H \subseteq S_M$
2. $\forall e^+ \in E^+ \exists A \in \text{AS}(B \cup H)$ such that A extends e^+

A hypothesis H is a violating solution if and only if

1. $H \subseteq S_M$
2. $\forall e^+ \in E^+ \exists A \in \text{Answer Sets}(B \cup H)$ such that A extends e^+
3. $\exists e^- \in E^- \exists A \in \text{Answer Sets}(B \cup H)$ such that A extends e^-

Given both definitions of positive and violating solutions, ILP_{LAS} is positive solutions that are not violating solutions.

Learning from Ordered Answer Sets (LOAS)

LOAS is an extension of ILASP by incorporating learning capability of weak constraints, is useful for modelling the agent's preferences ([13]). Examples used for learning are ordered pairs of partial answer sets, which are of the form $o = \langle e_1, e_2 \rangle$. An ASP Program P *bravely respects* o if and only if

1. $\exists A_1, A_2 \in \text{Answer Sets}(P)$ such that A_1 extends e_1 , A_2 extends e_2 , and $A_1 \succ_P A_2$

P *cautiously respects* o if and only if

1. $\forall A_1, A_2 \in \text{Answer Sets}(P)$ such that A_1 extends e_1 , A_2 extends e_2 , and $A_1 \succ_P A_2$

LOAS task is of the form $T = \langle B, S_M, E_+, E_-, O^b, O^c \rangle$ where O^b and O^c are brave and cautious orderings respectively, which are sets of ordering examples over set of positive partial interpretations E^+ . A hypothesis H is an inductive solution of T if and only if

1. $H \subseteq S_M$ in $ILP_{LOAS}(T)$
2. $H' \in ILP_{LAS}(\langle B, S_{LAS}(M_h, M_b), E^+, E^- \rangle)$ where H' is the subset of H without weak constraints
3. $\forall o \in O^b B \cup H$ bravely respects o
4. $\forall o \in O^c B \cup H$ cautiously respects o

A Context-dependent Learning from Ordered Answer Sets

Context-dependent learning from ordered answer sets ($ILP_{LOAS}^{context}$) is a further generalisation of ILP_{LOAS} with *context-dependent examples*. Context-dependent examples are examples to which each unique background knowledge (context) only applies. This way the background knowledge is more structured rather than one fixed background knowledge that are applied to all examples. Formally, partial interpretation is of the form $\langle e, C \rangle$ (called *context-dependent partial interpretation (CDPI)*), where e is a partial interpretation and C is called *context*, or an ASP program without weak constraints. A *context-dependent ordering example (CDOE)* is of the form $\langle \langle e_1, C_1 \rangle, \langle e_2, C_2 \rangle \rangle$. An APS program P *bravely respects* o if and only if

1. $\exists \langle A_1, A_2 \rangle$ such that $A_1 \in \text{Answer Sets}(P \cup C_1)$, $A_2 \in \text{Answer Sets}(P \cup C_2)$, A_1 extends e_1 , A_2 extends e_2 and $A_1 \prec_P A_2$

Similarly, an APS program P *cautiously respects* o if and only if

1. $\forall \langle A_1, A_2 \rangle$ such that $A_1 \in \text{Answer Sets}(P \cup C_1)$, $A_2 \in \text{Answer Sets}(P \cup C_2)$, A_1 extends e_1 , A_2 extends e_2 and $A_1 \prec_P A_2$

$ILP_{LOAS}^{context}$ task is of the form $T = \langle B, S_M, E_+, E_-, O^b, O^c \rangle$ where O^b and O^c are brave and cautious orderings respectively, which are sets of ordering examples over set of positive partial interpretations E^+ . A hypothesis H is an inductive solution of T if and only if

1. $H \subseteq S_M$ in $ILP_{LOAS}^{context}$
2. $\forall \langle e, C \rangle \in E^+, \exists A \exists A \in \text{Answer Sets}(B \cup C \cup H)$ such that A extends e
3. $\forall \langle e, C \rangle \in E^-, \nexists A \exists A \in \text{Answer Sets}(B \cup C \cup H)$ such that A extends e
4. $\forall o \in O^b B \cup H$ bravely respects o
5. $\forall o \in O^c B \cup H$ cautiously respects o

Two advantages of adding context-dependent are it increases the efficiency of learning tasks, and more expressive structure of the background knowledge to particular examples.

2.2 Reinforcement Learning (RL)

Reinforcement learning (RL) is another subfield of machine learning regarding an agent's behaviour in an environment in order to maximise the total reward. As shown in Figure 2.1, the agent interacts with an environment, and at each time step, an agent takes an action and receive observation, which affects the environment state and the reward (or penalty) it receives by the action outcome. In this section, we briefly introduce necessary background in RL for our research.




Figure 2.1: Agent and Environment

2.2.1 Markov Decision Process (MDP)

An agent interacts with an environment at a sequence of discrete time step, which is a part of the sequential history of observations, actions and rewards. The sequential history is formalised as $H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t$. A *state* is a function of the history $S_t = f(H_t)$, which determines the next environment. A state S_t is said to have Markov property if and only if $P[S_{t+1} | S_t] = P[S_{t+1} | S_1, \dots, S_t]$. In other words, the probability of reaching S_{t+1} depends only on S_t , which captures all the relevant information from earlier history ([19]).

When an agent must make a sequence of decision, the sequential decision problem can be formalised using Markov decision process (MDP). MDPs formally represent a fully observable environment of an agent for RL.

A MDP is of the form $\langle S, A, T_a, R_a, \gamma \rangle$ where:

- S is the set of finite states that is observable in the environment
- A is the set of finite actions taken by the agent
- $T_a(s, s')$ is a state transition in the form of probability matrix $\Pr(S_{t+1} = s' | s_t = s, a_t = a)$, which is the probability that action a in state s at time t will result in state s' at time $t+1$.

- R is a reward function $R_a(s, s') = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$, the expected immediate reward that action a in state s at time t will return
- γ is a discount factor $\gamma \in [0,1]$, which represents the preference of the agent for present rewards over future rewards

2.2.2 Policies and Value Functions

value functions estimate the expected return, or expected future rewarded, for a given action in a given state. The expected reward for an agent is dependent on the agent's action. A solution to the sequential decision problem is called a *policy* π , a sequence of actions that leads to a solution.

The state value function $v_\pi(s)$ of an MDP under a policy π is the expected return starting from state s , which is of the form:

$$v_\pi(s) = \mathbb{E}[G_t \mid S_t = s]$$

where $G_t = R_{t+1} + \gamma R_{t+2} + \dots$ (the total discounted reward from t).

A policy π is optimal if it maximises the action-value function. The transition and reward functions are not necessary to be known to compute π .

Among all possible value functions, an *optimal policy* π^* is the one that maximise the total rewards in the environment.

$$v^*(s) = \max_{\pi} v_\pi(s)$$

The optimal policy π^* corresponds to $v^*(s)$

$$\pi^* = \arg \max_{\pi} v_\pi(s)$$

TODO on-policy and off-policy learning.

$$A_t = \pi(S_t)$$

One of the common possibilities is that the agent chooses an action in order to maximise the discounted sum of future rewards.

$$A_t \text{ to maximise } R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$$

An action-value function evaluate a particular state by taking an action according the policy

$$q_\pi = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s, A_t = a, A_{t+1} = a,]$$

2.2.3 Model-based vs Model-free Learning

A model M is a representation of an MDP where the state space and action space are assumed to be known. The model represents state transitions and rewards. In this case, the problem to be solved becomes a planning problem for a series of actions to archive the agent's goal. Most of the reinforcement learning problems are model-free learning, where M is not unknown and the agent learns to achieve the goal by solely interacting with the environment. Thus the agent knows only possible states and actions, and the transition state and reward probability functions are unknown. The performance of model-based RL is limited to optimal policy given the model M . In other words, when the model is not a representation of the true MDP, the planning algorithms will not lead to the optimal policy, but suboptimal policy.

One algorithm which combine both aspects of model-based and model-free learning to solve the issue of sub-optimality is called Dyna ([25]), which is shown in Figure 2.2.




Figure 2.2: Relationships among learning, planning and acting ??

Dyna learns a model from real experience and use the model to generate simulated experience to update the evaluation functions.

This approach is more effective because the simulated experience is relatively easy to generate compared to having to have real experience, thus less iterations.

2.2.4 Temporal-Difference (TD) Learning

To solve a MDP, one of the approaches is called Temporal-Difference (TD) Learning. TD is an online model-free learning and learns directly from episodes of incomplete experiences without a model of the environment. TD updates the estimate by using the estimates of value function by bootstrap, which is formalised as

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)] \quad (2.2)$$

where $R_{t+1} + \gamma V(S_{t+1})$ is the target for TD update, which is biased estimated of $v_\pi(S_t)$, and $\delta = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called TD error, which is the error in $V(S_t)$ available at time $t+1$.

Since TD methods only needs to know the estimate of one step ahead and does not need the final outcome of the episodes, it can learn online after every time step. TD also works without the terminal state, which is the goal for an agent.

TD(0) is proved to converge to v_π in the table-based case (non-function approximation). However, because bootstrapping is update an estimate for an estimate, some bias is inevitable. And TD method is sensitive to initial value, (but low variance).

2.2.5 Q-Learning

Q-learning is off-policy TD learning defined in [26], where the agent only knows about the possible states and actions. The transition states and reward probability

functions are unknown to the agent. It is of the form:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(R_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)) \quad (2.3)$$

where α is the learning rate, γ is a discount rate between 0 and 1. The equation is used to update the state-action value function called Q function. The function $Q(S,A)$ predicts the best action A in state S to maximise the total cumulative rewards. The optimal Q-function $Q^*(s,a)$ is directly approximated by the learned action-value function Q.

Q-learning learns the value of its deterministic greedy policy from the experience and gradually converge to the optimal Q-function. It also explored following ϵ -greedy policy, which is a stochastic greedy policy, but with the probability of ϵ , the agent chooses an action randomly instead of the greedy action.

Chapter 3

Related Work

In this section, I summarise recent studies related to symbolic (deep) reinforcement learning.

?? introduced Deep Symbolic Reinforcement Learning (DSRL), a proof of concept for incorporating symbolic front end as a means of converting low-dimensional symbolic representation into spatio-temporal representations, which will be the state transitions input of reinforcement learning. DSRL extracts features using convolutional neural networks (CNNs) [27] and autoencoder, which are transformed into symbolic representations for relevant object types and positions of the objects, which represents abstract state-space. These abstract representations are the input for Q-learning algorithm to learn a policy on this particular state-space. DSRL was shown to outperform DRL on a stochastic variant environments. However, there are a number of drawbacks on this approach. First, the extraction of the individual objects was done by manually defined threshold of feature activation values, given that the games were geometrically simple. Thus this approach would not scale in a slightly complex game scenarios. Second, using deep neural network front-end might also cause a problem. As demonstrated in [24], a single irrelevant pixel could dramatically influence state through the change in CNNs. Third, the proposed method successfully used symbolic representations to achieve more data-efficient learning. However, there is still potential to apply symbolic learning to those symbolic representations to further improve the learning efficiency, which is what we attempt to do in this paper. ?? further explored this symbolic abstraction approach by incorporating relative position of each object with respect to every other object rather than absolute object position. They also assign priority to each Q-value function based on the relative distance of objects from an agent.

[28] added relational reinforcement learning, a classical subfield of research aiming to combine reinforcement learning with relational learning or Inductive Logic Programming, on top of DSRL approach added more abstract planning, and applied to much more complicated game environment than what was used in ??.

This idea of adding planning capability is aligned to our approach of using ILP to improve a RL agent. In this paper, we explore how to effectively learn the model of the environment and effectively use it to facilitate data-efficient learning and transfer learning capability.

Another approach for using symbolic reinforcement learning is storing heuristics

expressed by knowledge bases [[1]]. An agent learns the concept of *Hierarchical Knowledge Bases (HKBs)* [which is defined in more details in [3] and [2]] at every iteration of training, which contains multiple rules (state-action pairs). The agent then is able to decide itself when it should exploit the heuristic rather than the state-action pairs of the RL using *Strategic Depth*. This approach effectively used the heuristic knowledge bases, which acts as a sym-symbolic model of the game.

Another related field to our research is combining ASP and RL. The original concept of combining ASP and RL was in [7], where they developed an algorithm that efficiently finds the optimal solution of an MDP of non-stationary domains by using ASP to find the possible trajectories of an MDP. This approach focused more on efficient update on Q function rather than inductive learning. In order to find stationary sets, the extension of ASP called BC^+ , an action language, was used which can be directly translated the agent's actions into ASP form, and provide sequences of actions in Answer Sets.

ASP BC^+ translation was done manually.

Chapter 4

Project Overview

4.1 Proposed Architecture

The proposed tentative architecture is shown in Figure 4.1. The overall architecture is based on Sutton’s Dyna architecture ([25]), which combined both model-free learning and model-based learning. What is different from Dyna is the way we generate the model of the environment, which is explained in the following subsections.

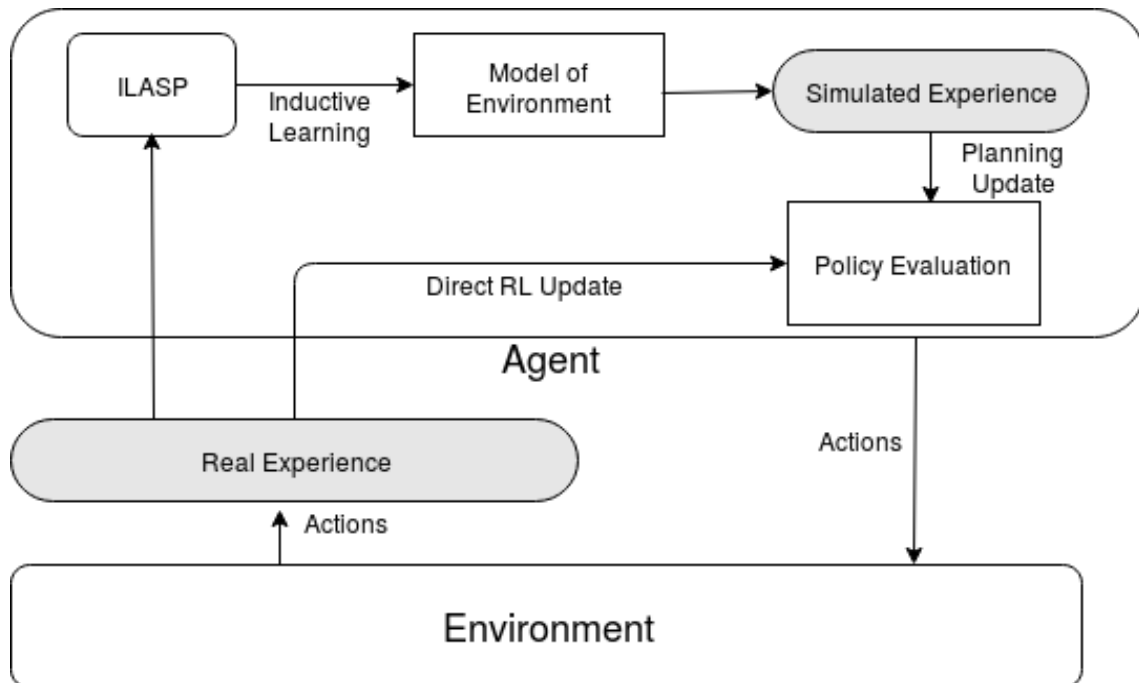


Figure 4.1: Proposed reinforcement learning architecture. ILASP learns to generate a model and updates based on the interaction with the environment, which is used to facilitate the policy evaluation.

4.1.1 ASP Representation Generation

At the first stage, the input of real experience needs to be converted in ASP form, which can be used to proceed the inductive learning in ILASP. The input used in ILASP is state transition, reward and action of the agent, which can be directly converted using a simple mapping table or an action language (such as BC^+ as used in [7]). The following ASP input is what is sent to ILASP.

```
agent_at_before(cell(X,Y), T).
agent_at_after(cell(X,Y), T).
reward(R, T).
action(A, T).
```

4.1.2 Model Generation and Update using ILASP

Once the input of the real experience is converted in ASP syntax, the agent should learn the following definition of the model of the environment using ILASP.

```
valid_move(C, T):- link(C, T).
link(C2, T):- agent_at(C1, T), adjacent(C0, C2), not obstacle(C0, C2).
agent_at(C, T):- agent_at_after(C, T)
```

The background knowledge is empty, and there are only positive examples in learning this task. Each example contains different transition history of the agent. Inclusions are valid moves and exclusions are invalid moves. Learning `valid_move` is the same as learning the rule of the games (the model of the environment), and it is updated as the agent explores in the environment.

In addition to the rule of the game, learning the following concepts will be crucial for transfer learning, as these concept will be applicable to any types of game environments.

```
adjacent(C1, C2):- cell(C1), cell(C2).
obstacle(C1, C2):- wall(C1, C2).
wall(C1, C2):- agent_at_before(C1, T1), agent_at_after(C1, T2)
enemy(C1, C2):- agent_at_before(C1, T1), agent_at_after(C2, T2), reward(R), R < -100
```

where it is assumed that the reward of -100 means losing the game or losing the agent's life. Once the agent has learned the concept of the game, it knows how to avoid the an obstacle in the adjacent cell in a new environment. Figure 4.2 illustrates this transfer capability.

If the agent A learnt the concept *adjacent* in Environment 1, the agent could reason (have a policy), for example, "if the adjacent cell is a wall, moving into that cell is an invalid move", "if the adjacent cell is an enemy, moving into that cell will incur negative reward". These learnt policies could immediately used in a new environment (e.g Environment 2), even though the agent's coordinates position is different




Figure 4.2: Simple Grid World to highlight transfer capability

from that in previous environment (e.g Environment 1).

4.1.3 Generation of Simulated Experience

Assuming the model is true representation of the partial environment, we can generate simulated experience to update Q function. The simulated experience is cheaper than real experience, which should contribute to efficient learning. However Dyna architecture is known to lead to sub-optimal policy when the model is not accurate, and the solution to this problem needs to be further considered.

Alternative to using simulated experience is to directly update state representations. Once the concept "adjacent" was learnt from a real experience, it could be incorporated in a state.

In addition, when the environment changes during the exploration, the agent should generalise its internal model to cope with the new environment in two different but similar environments, the agent's model is more general which covers both games. Thus in theory, the agent should still be able to achieve more efficient learning from the model-based learning, facilitating the transfer learning. Implementation of exactly how the model is generalised will be further investigated.

4.2 Further Research and Plan

The next phase of the research is further research on the proposed architecture and implementation and experiments.

4.2.1 Further Research

The proposed architecture is not finalised and will be reviewed regularly as the research is proceeded. More research needs to be devoted in finalising the overall architecture, and especially the following issues needs to be considered.

- Further investigation of whether ILASP can learn the concept of adjacent, which is crucial concept to know in any environment.

- How to generalise the agent's model when the environment changes. The new environment could be very similar to the previous one, or could be a completely different environment thus the agent should create a new internal model rather than generalising the existing model.
- The current proposed architecture is based on Dyna with simulated experiences. However, this might not be the best overall architecture, and the feasibility of using simulated experience with the learnt model with ILASP needs to be further investigated.
- Possibility of using other representational concepts such as *Predictive Representations of State* or *Affordance* [23] for the agent's learning task. These concept have not been considered at the moment, but could help better transfer learning.
- Preparation for a backup plan in case ILASP approach does not work, so that the reseach is feasible within 3 months of the reseach period.

4.2.2 Experiment Implementation

Once the architecture is decided, we will implement it using Python and ILASPv3.1.0 (Clingo) [14], and compare the performance with other learning methods using a game platform called GVG-AI Framework.

Other learning methods to be compared as benchmark will be Q-Learning and Deep Q-Learning since these methods are widely used in other related works. Other symbolic-based approaches could also be compared as an extension. The two main measurements for the performance of our new architecture are learning efficiency and transfer Learning capability as we stated our motivation in the Introduction.

GVG-AI Framework was created for the General Video Gamea AI Competition ¹, a game environment for an agent that should be able to play a wide variety of games without knowing which games are to be played. The underline language is the Video Game Definition Language (VGDL), which is a high-level description language for 2D video games providing a platform for computational intelligence research ([21]).

¹<http://www.gvgai.net/>

Chapter 5

Ethics Checklist

	Yes	No
Section 1: HUMAN EMBRYOS/FOETUSES		
Does your project involve Human Embryonic Stem Cells?		✓
Does your project involve the use of human embryos?		✓
Does your project involve the use of human foetal tissues / cells?		✓
Section 2: HUMANS		
Does your project involve human participants?		✓
Section 3: HUMAN CELLS / TISSUES		
Does your project involve human cells or tissues? (Other than from Human Embryos/Foetuses i.e. Section 1)?		✓
Section 4: PROTECTION OF PERSONAL DATA		
Does your project involve personal data collection and/or processing?		✓
Does it involve the collection and/or processing of sensitive personal data (e.g. health, sexual lifestyle, ethnicity, political opinion, religious or philosophical conviction)?		✓
Does it involve processing of genetic information?		✓
Does it involve tracking or observation of participants? It should be noted that this issue is not limited to surveillance or localization data. It also applies to Wan data such as IP address, MACs, cookies etc.		✓
Does your project involve further processing of previously collected personal data (secondary use)? For example Does your project involve merging existing data sets?		✓
Section 5: ANIMALS		
Does your project involve animals?		✓
Section 6: DEVELOPING COUNTRIES		
Does your project involve developing countries?		✓

If your project involves low and/or lower-middle income countries, are any benefit-sharing actions planned?		✓
Could the situation in the country put the individuals taking part in the project at risk?		✓
Section 7: ENVIRONMENTAL PROTECTION AND SAFETY		
Does your project involve the use of elements that may cause harm to the environment, animals or plants?		✓
Does your project deal with endangered fauna and/or flora /protected areas?		✓
Does your project involve the use of elements that may cause harm to humans, including project staff?		✓
Does your project involve other harmful materials or equipment, e.g. high-powered laser systems?		✓
Section 8: DUAL USE		
Does your project have the potential for military applications?		✓
Does your project have an exclusive civilian application focus?		✓
Will your project use or produce goods or information that will require export licenses in accordance with legislation on dual use items?		✓
Does your project affect current standards in military ethics e.g., global ban on weapons of mass destruction, issues of proportionality, discrimination of combatants and accountability in drone and autonomous robotics developments, incendiary or laser weapons?		✓
Section 9: MISUSE		
Does your project have the potential for malevolent/criminal/terrorist abuse?		✓
Does your project involve information on/or the use of biological-, chemical-, nuclear/radiological-security sensitive materials and explosives, and means of their delivery?		✓
Does your project involve the development of technologies or the creation of information that could have severe negative impacts on human rights standards (e.g. privacy, stigmatization, discrimination), if misapplied?	✓	
Does your project have the potential for terrorist or criminal abuse e.g. infrastructural vulnerability studies, cybersecurity related project?		✓
Section 10: LEGAL ISSUES		
Will your project use or produce software for which there are copyright licensing implications?	✓	

Will your project use or produce goods or information for which there are data protection, or other legal implications?		✓
Section 11: OTHER ETHICS ISSUES		
Are there any other ethics issues that should be taken into consideration?		✓

Bibliography

- [1] Daan Apeldoorn and Kern-Isberner Gabriele. An Agent-Based Learning Approach for Finding and Exploiting Heuristics in Unknown Environments. *Proceedings of the Thirteenth International Symposium on Commonsense Reasoning*, pages 1–8, 2017. pages 16
- [2] Daan Apeldoorn and Gabriele Kern-isberner. Towards an Understanding of What Is Learned : Extracting Multi-Abstraction-Level Knowledge from Learning Agents. pages 764–767. pages 16
- [3] Daan Apeldoorn and Gabriele Kern-isberner. When Should Learning Agents Switch to Explicit Knowledge ? 41:174–186, 2016. pages 16
- [4] Domenico Corapi, Alessandra Russo, and Emil Lupu. Inductive Logic Programming in Answer Set Programming. *Inductive Logic Programming*, pages 91–97, 2012. pages 7
- [5] Luc De Raedt. Logical settings for concept-learning. *Artificial Intelligence*, 95(1):187–201, 1997. pages 4, 7
- [6] Leonardo A. Ferreira, Reinaldo A. C. Bianchi, Paulo E. Santos, and Ramon Lopez de Mantaras. A method for the online construction of the set of states of a Markov Decision Process using Answer Set Programming. jun 2017. pages
- [7] Leonardo A. Ferreira, Reinaldo A. C. Bianchi, Paulo E. Santos, and Ramon Lopez de Mantaras. Answer Set Programming for Non-Stationary Markov Decision Processes. may 2017. pages 16, 18
- [8] Marta Garnelo, Kai Arulkumaran, and Murray Shanahan. Towards Deep Symbolic Reinforcement Learning. sep 2016. pages 2
- [9] Martin Gebser, Benjamin Kaufmann, Roland Kaminski, Max Ostrowski, Torsten Schaub, and Marius Schneider. Potassco: The Potsdam answer set solving collection. *AI Communications*, 24(2):107–124, 2011. pages 7
- [10] Michael Gelfond and Vladimir Lifschitz. The stable model semantics for logic programming. *5th International Conf. of Symp. on Logic Programming*, (December 2014):1070–1080, 1988. pages 6

- [11] Katsumi Inoue, Tony Ribeiro, and Chiaki Sakama. Learning from interpretation transition. *Machine Learning*, 94(1):51–79, 2014. pages 7
- [12] Mark Law, Alessandra Russo, and Krysia Broda. Inductive Learning of Answer Set Programs. *European Conference on Logics in Artificial Intelligence (JELIA)*, 2(Ray 2009):311–325, 2014. pages 8
- [13] Mark Law, Alessandra Russo, and Krysia Broda. Learning Weak Constraints in Answer Set Programming. jul 2015. pages 9
- [14] Mark Law, Alessandra Russo, and Krysia Broda. Inductive Learning of Answer Set Programs v3.1.0 User Manual. 2017. pages 20
- [15] Sergey Levine, Chelsea Finn, Trevor Darrell, and Pieter Abbeel. End-to-End Training of Deep Visuomotor Policies. apr 2015. pages 2
- [16] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dhharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning. *Nature*, 2015. pages 2
- [17] S Muggleton. Inductive logic programming. *New Generation Computing*, 8(4):295–318, 1991. pages 4
- [18] Ramon P Otero. Induction of Stable Models. *Conference on Inductive Logic Programming (ILP)*, pages 193–205, 2001. pages 7
- [19] Puterman M.L. *Markov Decision Processes Discrete Stochastic Dynamic Programming*. John Wiley & Sons, Inc., New York, NY, 1994. pages 11
- [20] Chiaki Sakama and Katsumi Inoue. Brave induction: A logical framework for learning from incomplete information. *Machine Learning*, 76(1):3–35, 2009. pages 7
- [21] Tom Schaul. A video game description language for model-based or interactive learning, 2013. pages 20
- [22] David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George Van Den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529(7587):484–489, 2016. pages 2
- [23] Mohan Sridharan, Ben Meadows, and Rocio Gomez. What can I not do? Towards an Architecture for Reasoning about and Learning Affordances. *International Conference on Automated Planning and Scheduling*, (Icaps):461–469, 2017. pages 20

-
- [24] Jiawei Su, Danilo Vasconcellos Vargas, and Sakurai Kouichi. One pixel attack for fooling deep neural networks. oct 2017. pages 15
 - [25] Richard S. Sutton. Integrated Architectures for Learning, Planning, and Reacting Based on Approximating Dynamic Programming. *Machine Learning Proceedings 1990*, 02254(1987):216–224, 1990. pages 13, 17
 - [26] Watkins. Learning from Delayed Reward. pages 13
 - [27] Y.LeCun, L.Bottou, Y.Bengio and P.Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86:2278–2324, 1998. pages 15
 - [28] Vinicius Zambaldi, David Raposo, Adam Santoro, Victor Bapst, Yujia Li, Igor Babuschkin, Karl Tuyls, David Reichert, Timothy Lillicrap, Edward Lockhart, Murray Shanahan, Victoria Langston, Razvan Pascanu, Matthew Botvinick, Oriol Vinyals, and Peter Battaglia. Relational Deep Reinforcement Learning. jun 2018. pages 15